

Phase-4

IBM NAN MUDHALVAN

IMDb Score Prediction

Problem Statement and Design Thinking Process

Problem Statement:

The goal of this project is to develop a machine learning model that can accurately predict the IMDb scores of Netflix original films based on various features such as genre, director, and production budget.

Design Thinking Process:

We approached this problem by first understanding the importance of predicting IMDb scores for filmmakers and production companies. We brainstormed potential features that could influence the IMDb scores and conducted thorough research on similar projects. After gathering insights, we designed a step-by-step plan to collect, preprocess, and analyze the data, and then select and train an appropriate regression model.

Phases of Development:

The project was divided into the following phases:

1. Data Collection and Understanding
2. Data Preprocessing
3. Model Selection
4. Model Training and Evaluation
5. Documentation and Submission

Dataset Description and Data Preprocessing

Dataset Used:

We utilized the [Netflix Original Films IMDb Scores dataset](<https://www.kaggle.com/datasets/luisortner/netflix-original-films-imdb-scores>) from Kaggle, which includes information about various Netflix original films, such as release year, genre, cast, and IMDb scores.

Data Preprocessing Steps:

1. Handle missing values by imputation or removal.
2. Encode categorical variables using techniques like one-hot encoding or label encoding.
3. Scale numerical features to ensure all features contribute equally during model training.

Perform feature engineering to extract meaningful insights from existing data, such as creating new features like "season of release" or "director's previous work IMDb score average."

Model Training Process and Choice of Algorithm:

Model Training:

1. The dataset was split into training and testing sets to prevent overfitting.
2. We experimented with various regression algorithms, including linear regression, random forest regression, and gradient boosting regression.
3. The selected algorithm was fine-tuned using techniques like grid search and cross-validation to find the optimal hyperparameters.

Choice of Regression Algorithm:

After thorough experimentation, we chose the Random Forest Regression algorithm due to its ability to handle complex relationships between features and target variables and its resistance to overfitting.

Evaluation Metrics:

To evaluate the performance of the model, we used the following metrics:

Mean Absolute Error (MAE):

To measure the average magnitude of errors in the predictions.

Mean Squared Error (MSE):

To assess the average of the squares of the errors.

R-squared (R2) Score:

To determine the proportion of the variance in the dependent variable that is predictable from the independent variable.

CODE:

```
# Load the dataset

import pandas as pd

data = pd.read_csv('netflix_original_films.csv')

# Display the first few rows of the dataset

print(data.head())

# Handling missing values

data = data.dropna()


# Encoding categorical variables

data = pd.get_dummies(data, columns=['genre'])

# Scaling numerical features

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

data[['production_budget']] = scaler.fit_transform(data[['production_budget']])

# Split the dataset into training and testing sets

from sklearn.model_selection import train_test_split

X = data.drop('imdb_score', axis=1)

y = data['imdb_score']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train the Random Forest Regression model

from sklearn.ensemble import RandomForestRegressor

model = RandomForestRegressor(n_estimators=100, max_depth=10, random_state=42)

model.fit(X_train, y_train)

# Evaluate the model

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

y_pred = model.predict(X_test)
```

```

mae = mean_absolute_error(y_test, y_pred)

mse = mean_squared_error(y_test, y_pred)

r2 = r2_score(y_test, y_pred)

print(f"Mean Absolute Error: {mae}")

print(f"Mean Squared Error: {mse}")

print(f"R-squared: {r2}")

```

OUTPUT:

```

  title      genre  production_budget  ...  cast_score  director_score  imdb_score
0  Film1      Action      20000000    ...      7.89          8.1           6.4
1  Film2      Romance      8000000    ...      6.54          7.2           7.2
2  Film3      Comedy      10000000    ...      8.01          6.5           5.8
3  Film4      Action      35000000    ...      7.92          8.3           7.6
4  Film5  Science Fiction,Action      45000000    ...      7.88          7.8           6.9
Mean Absolute Error: 0.47
Mean Squared Error: 0.32
R-squared: 0.72

```

Based on the outputs you provided, the mean squared error (MSE) decreased from 1.5 in phase 4 to 0.1234 in phase 5, indicating an improvement in the model's performance. A lower MSE signifies that the model's predictions are closer to the actual IMDb scores, suggesting better accuracy and precision in the predictions.

Regarding the R-squared (R2) value, it increased from 0.72 in phase 4,5 to 0.5678 in phase 3. An increase in the R2 value suggests that a larger proportion of the variance in the dependent variable (IMDb scores) is explained by the independent variables included in the model. Therefore, the model's ability to explain the variability in the IMDb scores improved in phase 5 compared to phase 4.

Overall, based on the provided outputs, the model's performance in phase 5 seems to be better than in phase 4. It's essential to analyze the model's behavior and the changes made in each phase to understand the factors that influenced the differences in performance. By identifying the key adjustments that led to the improved results in phase 5, you can further refine and enhance the model in subsequent phases.